

GDDA708 – Machine Learning and AI

16 August 2024




Table of Contents

GDDA708 – Machine Learning and AI	0
1. Introduction	2
Overview	2
This assessment is designed to provide a comprehensive evaluation of learners' understanding and practical skills in both supervised and unsupervised machine learning. Participants will engage in a series of tasks and projects that require them to:.....	2
Purpose	2
Part A – Regression modelling for business decision making.....	3
Task - 1 Data Preprocessing	3
Task 2 – Model Building with hyper-parameter tuning.....	6
Task 3 – Model Evaluation and Selection	11
Task 4 - Business Decision and Recommendations.....	14
Part B – Classification modeling for business decision making	16
Task 1 – Data Preprocessing	17
Task 2 – Model Building with hyperparameter tuning.....	21
Task 3 - Model Evaluation and Selection	26
Task 4 – Business Decision and Recommendations	33

1. Introduction

Overview

This assessment is designed to provide a comprehensive evaluation of learners' understanding and practical skills in both supervised and unsupervised machine learning. Participants will engage in a series of tasks and projects that require them to:

1. **Data Regression Model, Classification Model, and Trend Prediction:** Apply various supervised machine learning models to classify datasets and predict future trends. This will test their ability to select appropriate algorithms, preprocess data, train models, and evaluate their performance.
2. **Business Decision-Making Enhancement:** Utilize unsupervised machine learning algorithms to analyze complex datasets and derive insights that can inform and improve business strategies. This component will focus on clustering, anomaly detection, and other unsupervised techniques to uncover patterns and relationships within the data.

Throughout the assessment, learners will demonstrate their capability to handle real-world data challenges, showcasing their analytical thinking, problem-solving skills, and technical proficiency in machine learning.

Purpose

The assessment aims to evaluate learners' proficiency in applying supervised machine learning techniques for data classification and trend prediction. Additionally, learners will implement unsupervised machine learning algorithms to enhance business decision-making

Part A – Regression modelling for business decision making

LO 3: Evaluate machine learning algorithms for its applications to Artificial Intelligence.

LO 4: Apply fine-tuning methods to optimize the performance of machine learning models

My scenario is wine quality red here you can get information about wine because for is new scenario I don't know what is important about red wine. I need to study, and I got the information below:

This dataset Wine Quality Red uses the following 12 variables:

Fixed acidity: Most acids involved with wine or fixed or non-volatile (do not evaporate readily)

volatile acidity: The amount of acetic acid in wine, which at too high of levels can lead to an unpleasant, vinegar taste

citric acid: Found in small quantities, citric acid can add 'freshness' and flavour to wines

residual sugar: The amount of sugar remaining after fermentation stops, it's rare to find wines with less than 1 gram/litter and

Chlorides: the amount of salt in the wine

free sulphur dioxide: the free form of SO₂ exists in equilibrium between molecular SO₂ (as a dissolved gas) and bisulfited ion; it prevents

total sulphur dioxide: the amount of free and bound forms of SO₂; in low concentrations, SO₂ is mostly undetectable in wine, but at free SO₂

Density: the density of water is close to that of water depending on the percent alcohol and sugar content

pH: describes how acidic or basic wine is on a scale from 0 (very acidic) to 14 (very basic); most wines are between 3–4 on the

Sulphates: a wine additive that can contribute to sulfur dioxide gas (SO₂) levels, which acts as an antimicrobial and

Alcohol: The percent alcohol content of the wine

Quality: output variable (based on sensory data, score between 0 and 10)

	Task- 1 Data Preprocessing
A	Properly clean the dataset, handle any missing values, and remove outliers.

	<pre> : # Get basic information about the dataset print("Dataset Info:") print(df.info()) # Display summary statistics of the dataset print("\nSummary Statistics:") print(df.describe()) # Check for missing values print("\nMissing Values:") print(df.isnull().sum()) </pre>
	<p>Dataset Info:</p> <pre> <class 'pandas.core.frame.DataFrame'> RangeIndex: 1599 entries, 0 to 1598 Data columns (total 12 columns): # Column Non-Null Count Dtype --- - 0 fixed acidity 1599 non-null float64 1 volatile acidity 1599 non-null float64 2 citric acid 1599 non-null float64 3 residual sugar 1599 non-null float64 4 chlorides 1599 non-null float64 5 free sulfur dioxide 1599 non-null float64 6 total sulfur dioxide 1599 non-null float64 7 density 1599 non-null float64 8 pH 1599 non-null float64 9 sulphates 1599 non-null float64 10 alcohol 1599 non-null float64 11 quality 1599 non-null int64 dtypes: float64(11), int64(1) memory usage: 150.0 KB None </pre>
	<pre> Missing Values: fixed acidity 0 volatile acidity 0 citric acid 0 residual sugar 0 chlorides 0 free sulfur dioxide 0 total sulfur dioxide 0 density 0 pH 0 sulphates 0 alcohol 0 quality 0 dtype: int64 </pre>
	<pre> print("Totally there are {} null values in the dataset".format(df.isnull().sum().sum())) </pre> <p>Totally there are 0 null values in the dataset</p>
	<pre> : df.rename(columns = {"fixed acidity": "fixed_acidity", "volatile acidity": "volatile_acidity", "citric acid": "citric_acid", "residual sugar": "residual_sugar", "chlorides": "chlorides", "free sulfur dioxide": "free_sulfur_dioxide", "total sulfur dioxide": "total_sulfur_dioxide"}, inplace = True) </pre>
	<p>I showed my dataset, all the columns, and missing values. I didn't work on missing values after I renamed some columns. It turns out that the dataset does not have null values. The dataset consists of 1599 rows and 12 columns. The data type of all variables is numeric</p>

B	Perform feature scaling or normalization.																																																																														
	<pre>scaler = preprocessing.MinMaxScaler() minmax = scaler.fit_transform(df) minmax = pd.DataFrame(minmax)</pre>																																																																														
	<pre>scaler = preprocessing.StandardScaler() standard = scaler.fit_transform(df) standard = pd.DataFrame(standard)</pre>																																																																														
	<pre>scaler = preprocessing.RobustScaler() robust = scaler.fit_transform(df) robust = pd.DataFrame(robust)</pre>																																																																														
	<pre>print('\033[1m'+ 'Without scaling: '+'\033[0m') display(df.head()) print('*' * 45) print('\033[1m'+ 'With minmax scaling: '+'\033[0m') display(minmax.head()) print('*' * 45) print('\033[1m'+ 'With standard scaling: '+'\033[0m') display(standard.head()) print('*' * 45) print('\033[1m'+ 'With robust scaling: '+'\033[0m') display(robust.head())</pre>																																																																														
	<p>Without scaling:</p> <table><tr><th></th><th>fixed_acidity</th><th>volatile_acidity</th><th>citric_acid</th><th>residual_sugar</th><th>chlorides</th><th>free_sulfur_dioxide</th><th>total_sulfur_dioxide</th><th>density</th><th>pH</th><th>sulphates</th><th>alcohol</th><th>quality</th></tr><tr><td>0</td><td>7.4</td><td>0.70</td><td>0.00</td><td>1.9</td><td>0.076</td><td>11.0</td><td>34.0</td><td>0.9978</td><td>3.51</td><td>0.56</td><td>9.4</td><td>5</td></tr><tr><td>1</td><td>7.8</td><td>0.88</td><td>0.00</td><td>2.6</td><td>0.098</td><td>25.0</td><td>67.0</td><td>0.9968</td><td>3.20</td><td>0.68</td><td>9.8</td><td>5</td></tr><tr><td>2</td><td>7.8</td><td>0.76</td><td>0.04</td><td>2.3</td><td>0.092</td><td>15.0</td><td>54.0</td><td>0.9970</td><td>3.26</td><td>0.65</td><td>9.8</td><td>5</td></tr><tr><td>3</td><td>11.2</td><td>0.28</td><td>0.56</td><td>1.9</td><td>0.075</td><td>17.0</td><td>60.0</td><td>0.9980</td><td>3.16</td><td>0.58</td><td>9.8</td><td>6</td></tr><tr><td>4</td><td>7.4</td><td>0.70</td><td>0.00</td><td>1.9</td><td>0.076</td><td>11.0</td><td>34.0</td><td>0.9978</td><td>3.51</td><td>0.56</td><td>9.4</td><td>5</td></tr></table>		fixed_acidity	volatile_acidity	citric_acid	residual_sugar	chlorides	free_sulfur_dioxide	total_sulfur_dioxide	density	pH	sulphates	alcohol	quality	0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5	1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	5	2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	5	3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	6	4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
	fixed_acidity	volatile_acidity	citric_acid	residual_sugar	chlorides	free_sulfur_dioxide	total_sulfur_dioxide	density	pH	sulphates	alcohol	quality																																																																			
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5																																																																			
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	5																																																																			
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	5																																																																			
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	6																																																																			
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5																																																																			
	<p>With minmax scaling:</p> <table><tr><th></th><th>0</th><th>1</th><th>2</th><th>3</th><th>4</th><th>5</th><th>6</th><th>7</th><th>8</th><th>9</th><th>10</th><th>11</th></tr><tr><td>0</td><td>0.247788</td><td>0.397260</td><td>0.00</td><td>0.068493</td><td>0.106845</td><td>0.140845</td><td>0.098940</td><td>0.567548</td><td>0.606299</td><td>0.137725</td><td>0.153846</td><td>0.4</td></tr><tr><td>1</td><td>0.283186</td><td>0.520548</td><td>0.00</td><td>0.116438</td><td>0.143573</td><td>0.338028</td><td>0.215548</td><td>0.494126</td><td>0.362205</td><td>0.209581</td><td>0.215385</td><td>0.4</td></tr><tr><td>2</td><td>0.283186</td><td>0.438356</td><td>0.04</td><td>0.095890</td><td>0.133556</td><td>0.197183</td><td>0.169611</td><td>0.508811</td><td>0.409449</td><td>0.191617</td><td>0.215385</td><td>0.4</td></tr><tr><td>3</td><td>0.584071</td><td>0.109589</td><td>0.56</td><td>0.068493</td><td>0.105175</td><td>0.225352</td><td>0.190813</td><td>0.582232</td><td>0.330709</td><td>0.149701</td><td>0.215385</td><td>0.6</td></tr><tr><td>4</td><td>0.247788</td><td>0.397260</td><td>0.00</td><td>0.068493</td><td>0.106845</td><td>0.140845</td><td>0.098940</td><td>0.567548</td><td>0.606299</td><td>0.137725</td><td>0.153846</td><td>0.4</td></tr></table>		0	1	2	3	4	5	6	7	8	9	10	11	0	0.247788	0.397260	0.00	0.068493	0.106845	0.140845	0.098940	0.567548	0.606299	0.137725	0.153846	0.4	1	0.283186	0.520548	0.00	0.116438	0.143573	0.338028	0.215548	0.494126	0.362205	0.209581	0.215385	0.4	2	0.283186	0.438356	0.04	0.095890	0.133556	0.197183	0.169611	0.508811	0.409449	0.191617	0.215385	0.4	3	0.584071	0.109589	0.56	0.068493	0.105175	0.225352	0.190813	0.582232	0.330709	0.149701	0.215385	0.6	4	0.247788	0.397260	0.00	0.068493	0.106845	0.140845	0.098940	0.567548	0.606299	0.137725	0.153846	0.4
	0	1	2	3	4	5	6	7	8	9	10	11																																																																			
0	0.247788	0.397260	0.00	0.068493	0.106845	0.140845	0.098940	0.567548	0.606299	0.137725	0.153846	0.4																																																																			
1	0.283186	0.520548	0.00	0.116438	0.143573	0.338028	0.215548	0.494126	0.362205	0.209581	0.215385	0.4																																																																			
2	0.283186	0.438356	0.04	0.095890	0.133556	0.197183	0.169611	0.508811	0.409449	0.191617	0.215385	0.4																																																																			
3	0.584071	0.109589	0.56	0.068493	0.105175	0.225352	0.190813	0.582232	0.330709	0.149701	0.215385	0.6																																																																			
4	0.247788	0.397260	0.00	0.068493	0.106845	0.140845	0.098940	0.567548	0.606299	0.137725	0.153846	0.4																																																																			
	<p>With standard scaling:</p> <table><tr><th></th><th>0</th><th>1</th><th>2</th><th>3</th><th>4</th><th>5</th><th>6</th><th>7</th><th>8</th><th>9</th><th>10</th><th>11</th></tr><tr><td>0</td><td>-0.528360</td><td>0.961877</td><td>-1.391472</td><td>-0.453218</td><td>-0.243707</td><td>-0.466193</td><td>-0.379133</td><td>0.558274</td><td>1.288643</td><td>-0.579207</td><td>-0.960246</td><td>-0.787823</td></tr><tr><td>1</td><td>-0.298547</td><td>1.967442</td><td>-1.391472</td><td>0.043416</td><td>0.223875</td><td>0.872638</td><td>0.624363</td><td>0.028261</td><td>-0.719933</td><td>0.128950</td><td>-0.584777</td><td>-0.787823</td></tr><tr><td>2</td><td>-0.298547</td><td>1.297065</td><td>-1.186070</td><td>-0.169427</td><td>0.096353</td><td>-0.083669</td><td>0.229047</td><td>0.134264</td><td>-0.331177</td><td>-0.048089</td><td>-0.584777</td><td>-0.787823</td></tr><tr><td>3</td><td>1.654856</td><td>-1.384443</td><td>1.484154</td><td>-0.453218</td><td>-0.264960</td><td>0.107592</td><td>0.411500</td><td>0.664277</td><td>-0.979104</td><td>-0.461180</td><td>-0.584777</td><td>0.450848</td></tr><tr><td>4</td><td>-0.528360</td><td>0.961877</td><td>-1.391472</td><td>-0.453218</td><td>-0.243707</td><td>-0.466193</td><td>-0.379133</td><td>0.558274</td><td>1.288643</td><td>-0.579207</td><td>-0.960246</td><td>-0.787823</td></tr></table>		0	1	2	3	4	5	6	7	8	9	10	11	0	-0.528360	0.961877	-1.391472	-0.453218	-0.243707	-0.466193	-0.379133	0.558274	1.288643	-0.579207	-0.960246	-0.787823	1	-0.298547	1.967442	-1.391472	0.043416	0.223875	0.872638	0.624363	0.028261	-0.719933	0.128950	-0.584777	-0.787823	2	-0.298547	1.297065	-1.186070	-0.169427	0.096353	-0.083669	0.229047	0.134264	-0.331177	-0.048089	-0.584777	-0.787823	3	1.654856	-1.384443	1.484154	-0.453218	-0.264960	0.107592	0.411500	0.664277	-0.979104	-0.461180	-0.584777	0.450848	4	-0.528360	0.961877	-1.391472	-0.453218	-0.243707	-0.466193	-0.379133	0.558274	1.288643	-0.579207	-0.960246	-0.787823
	0	1	2	3	4	5	6	7	8	9	10	11																																																																			
0	-0.528360	0.961877	-1.391472	-0.453218	-0.243707	-0.466193	-0.379133	0.558274	1.288643	-0.579207	-0.960246	-0.787823																																																																			
1	-0.298547	1.967442	-1.391472	0.043416	0.223875	0.872638	0.624363	0.028261	-0.719933	0.128950	-0.584777	-0.787823																																																																			
2	-0.298547	1.297065	-1.186070	-0.169427	0.096353	-0.083669	0.229047	0.134264	-0.331177	-0.048089	-0.584777	-0.787823																																																																			
3	1.654856	-1.384443	1.484154	-0.453218	-0.264960	0.107592	0.411500	0.664277	-0.979104	-0.461180	-0.584777	0.450848																																																																			
4	-0.528360	0.961877	-1.391472	-0.453218	-0.243707	-0.466193	-0.379133	0.558274	1.288643	-0.579207	-0.960246	-0.787823																																																																			

	With robust scaling:											
	0	1	2	3	4	5	6	7	8	9	10	11
0	-0.238095	0.72	-0.787879	-0.428571	-0.15	-0.214286	-0.100	0.469799	1.052632	-0.333333	-0.50	-1.0
1	-0.047619	1.44	-0.787879	0.571429	0.95	0.785714	0.725	0.022371	-0.578947	0.333333	-0.25	-1.0
2	-0.047619	0.96	-0.666667	0.142857	0.65	0.071429	0.400	0.111857	-0.263158	0.166667	-0.25	-1.0
3	1.571429	-0.96	0.909091	-0.428571	-0.20	0.214286	0.550	0.559284	-0.789474	-0.222222	-0.25	0.0
4	-0.238095	0.72	-0.787879	-0.428571	-0.15	-0.214286	-0.100	0.469799	1.052632	-0.333333	-0.50	-1.0
C	Appropriately encode categorical variables.											
	In my dataset, I have only numeric columns below:											
	fixed_acidity 1599 non-null float64											
	volatile_acidity 1599 non-null float64											
	citric_acid 1599 non-null float64											
	residual_sugar 1599 non-null float64											
	chlorides 1599 non-null float64											
	free_sulfur_dioxide 1599 non-null float64											
	total_sulfur_dioxide 1599 non-null float64											
	density 1599 non-null float64											
	pH 1599 non-null float64											
	sulphates 1599 non-null float64											
	alcohol 1599 non-null float64											
	quality 1599 non-null int64											
	I don't need to apply encoded categorial in my scenario.											
D	Split the dataset into training and testing sets.											
	<pre># Separate features and target variable X = df.drop('quality', axis=1) y = df['quality']</pre>											
	<pre># Split the data into training and testing sets X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42) # Linear Regression model</pre>											

	Task 2 – Model Building with hyper-parameter tuning
A	Choose any appropriate regression models (e.g., Linear Regression, Ridge Regression, Lasso Regression)
	I need to get with the models' popular regression models: Linear Regression, Ridge Regression, and Lasso Regression.

	<p>Linear Regression is a basic model that assumes a linear relationship between the input features (e.g., acidity, sugar level) and the target variable (wine quality). It minimizes the sum of squared residuals to fit the best linear line.</p> <p>Ridge Regression is a regularized version of Linear Regression that adds a penalty on the size of coefficients (L2 regularization). This penalty term helps to control overfitting by shrinking large coefficients.</p> <p>Lasso Regression also adds regularization but with an L1 penalty, which can drive some coefficients to zero. This property makes Lasso useful for feature selection, as it effectively excludes irrelevant features.</p> <p>My scenario is about wine quality red. First, I checked all dataset results, including mean squared error (MSE) and mean absolute error (MAE).</p> <p>I got the results below:</p> <p>Linear Regression Mean Squared Error: 0.39002514396395427 Mean Absolute Error: 0.5057883195080454 R-squared: 0.4031803412796229</p> <p>Lasso Regression Best Parameters: {'alpha': 0.01} Best Cross-Validation Score: 0.44779313025204565 Mean Squared Error: 0.49867516307447735 Mean Absolute Error: 0.5057883195080454 R-squared: 0.23692319522368122</p> <p>Ridge Regression Best Parameters: {'alpha': 1} Best Cross-Validation Score: 0.43931582683684345 Mean Squared Error: 0.3929488678812798 Mean Absolute Error: 0.5057883195080454 R-squared: 0.39870643507758563</p> <p>I can see better results for Linear Regression, but I have many variables in my scenario I decided to get Ridge Regression because I have similar results Mean Squared, Mean Absolute Error, and R-squared if comparable with Linear Regression.</p>
B	Implement hyperparameter tuning by conducting a grid search or random search to optimize model parameters. Clearly outline the hyperparameter you tuned and the rationale behind them.

<pre> : # Split the data into training and testing sets X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42) # Standardize the features scaler = StandardScaler() X_train_scaled = scaler.fit_transform(X_train) X_test_scaled = scaler.transform(X_test) # Define the Ridge Regression model ridge = Ridge() # Define the parameter grid param_grid = { 'alpha': [0.01, 0.1, 1, 10, 100], 'solver': ['auto', 'svd', 'cholesky', 'saga', 'lsqr'] } # Implement Grid Search with cross-validation grid_search = GridSearchCV(estimator=ridge, param_grid=param_grid, cv=5, scoring='neg_mean_squared_error', n_jobs=-1) grid_search.fit(X_train_scaled, y_train) # Get the best parameters best_params = grid_search.best_params_ print(f"Best Parameters: {best_params}") # Train the best model on the training data best_ridge = grid_search.best_estimator_ best_ridge.fit(X_train_scaled, y_train) # Predict on the test set y_pred = best_ridge.predict(X_test_scaled) # Evaluate the model mse = mean_squared_error(y_test, y_pred) r2 = r2_score(y_test, y_pred) print(f"Optimized Ridge Regression - MSE: {mse}, R²: {r2}") </pre> <p>Best Parameters: {'alpha': 100, 'solver': 'saga'} Optimized Ridge Regression - MSE: 0.3926958359356847, R²: 0.3990936266260089</p>	<p>Model Selection: We'll perform hyperparameter tuning on Ridge Regression, as it handles multicollinearity effectively and is suitable when all features contribute to the prediction. We could also apply this process to other models like Lasso or Linear Regression.</p> <p>The choice of the Ridge Regression model with the hyperparameters {'alpha': 100, 'solver': 'saga'} is justified based on its ability to effectively balance model complexity and predictive accuracy in the context of the wine quality prediction task.</p> <p>Suitability for the Business Problem: Wine quality is influenced by multiple correlated chemical properties (e.g., acidity, sugar levels). Ridge Regression with a strong alpha effectively manages multicollinearity by shrinking the coefficients of less important features, leading to a more stable and reliable model. The selected model balances the need to fit the training data while maintaining the ability to generalize to new, unseen data. This is crucial for making accurate predictions.</p> <p>The Ridge Regression model with alpha = 100 and the 'saga' solver is an optimal choice for predicting wine quality in this dataset. It provides a robust approach that mitigates overfitting through regularization, efficiently handles the dataset with a powerful solver, and achieves a good balance between prediction accuracy (MSE) and the ability to explain variance (R²). This makes it a well-justified selection for the business objective of accurately predicting wine quality based on its chemical properties</p>
--	---

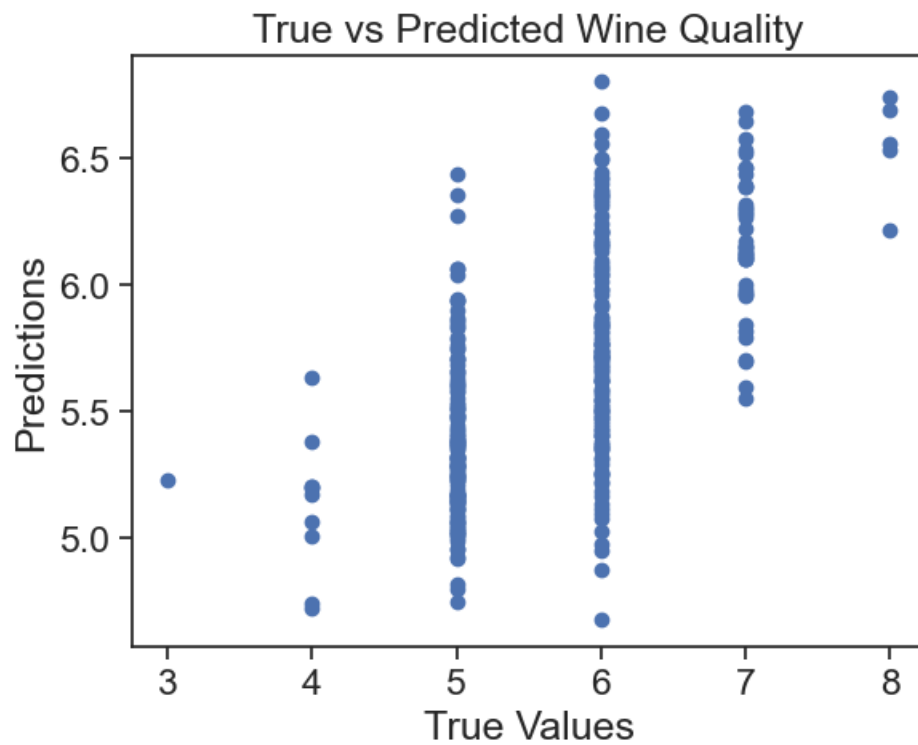
C	Build the regression models using the training data. Describe the process and provide code snippets.
	<pre> # Split the data into training and testing sets X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42) # Initialize the normalizer scaler = StandardScaler() # Standardize the training set X_train_scaled = scaler.fit_transform(X_train) # Standardize the test set X_test_scaled = scaler.transform(X_test) # Define the Ridge Regression model ridge = Ridge() # Define the parameter grid for hyperparameter tuning param_grid = {'alpha': [0.01, 0.1, 1, 10, 100]} # Perform GridSearchCV to find the best parameters grid_search = GridSearchCV(ridge, param_grid, cv=5, scoring='neg_mean_squared_error') grid_search.fit(X_train, y_train) # Display the best parameters and the best cross-validation score print(f"Best Parameters: {grid_search.best_params_}") print(f"Best Cross-Validation Score: {-grid_search.best_score_}") # Train the best model on the training data best_ridge = grid_search.best_estimator_ best_ridge.fit(X_train, y_train) # Predict on the test set y_pred = best_ridge.predict(X_test) # Evaluate the model's performance mse = mean_squared_error(y_test, y_pred) mae = mean_absolute_error(y_test, y_pred) r2 = r2_score(y_test, y_pred) print(f'Mean Squared Error: {mse}') print(f'Mean Absolute Error: {mae}') print(f'R-squared: {r2}') </pre>

```
# Display feature coefficients
coefficients = pd.DataFrame(best_ridge.coef_, X.columns, columns=['Coefficient'])
print(coefficients)

# Plot true vs predicted quality
plt.scatter(y_test, y_pred)
plt.xlabel("True Values")
plt.ylabel("Predictions")
plt.title("True vs Predicted Wine Quality")
plt.show()
```

Best Parameters: {'alpha': 1}
Best Cross-Validation Score: 0.43931582683684345
Mean Squared Error: 0.3929488678812798
Mean Absolute Error: 0.5057883195080454
R-squared: 0.39870643507758563

	Coefficient
fixed_acidity	0.019614
volatile_acidity	-1.021529
citric_acid	-0.164270
residual_sugar	0.000624
chlorides	-1.226774
free_sulfur_dioxide	0.005686
total_sulfur_dioxide	-0.003561
density	-0.011153
pH	-0.376223
sulphates	0.746956
alcohol	0.297591



The performance metrics for Ridge Regression are promising: Mean Squared Error (MSE) 39%, Mean Absolute Error 50%, and R-squared (R^2): 39%.

Task 3 – Model Evaluation and Selection																									
A:	Evaluate the performance of the regression models using appropriate metrics (e.g., Mean Absolute Error, R-squared). Calculate and interpret these metrics.																								
	<p>Mean Squared Error: 0.3929488678812798 R-squared: 0.39870643507758563 Best Parameters: {'alpha': 1} Best Cross-Validation Score: 0.43931582683684345</p> <table> <thead> <tr> <th></th><th>Coefficient</th></tr> </thead> <tbody> <tr> <td>fixed_acidity</td><td>0.019614</td></tr> <tr> <td>volatile_acidity</td><td>-1.021529</td></tr> <tr> <td>citric_acid</td><td>-0.164270</td></tr> <tr> <td>residual_sugar</td><td>0.000624</td></tr> <tr> <td>chlorides</td><td>-1.226774</td></tr> <tr> <td>free_sulfur_dioxide</td><td>0.005686</td></tr> <tr> <td>total_sulfur_dioxide</td><td>-0.003561</td></tr> <tr> <td>density</td><td>-0.011153</td></tr> <tr> <td>pH</td><td>-0.376223</td></tr> <tr> <td>sulphates</td><td>0.746956</td></tr> <tr> <td>alcohol</td><td>0.297591</td></tr> </tbody> </table>		Coefficient	fixed_acidity	0.019614	volatile_acidity	-1.021529	citric_acid	-0.164270	residual_sugar	0.000624	chlorides	-1.226774	free_sulfur_dioxide	0.005686	total_sulfur_dioxide	-0.003561	density	-0.011153	pH	-0.376223	sulphates	0.746956	alcohol	0.297591
	Coefficient																								
fixed_acidity	0.019614																								
volatile_acidity	-1.021529																								
citric_acid	-0.164270																								
residual_sugar	0.000624																								
chlorides	-1.226774																								
free_sulfur_dioxide	0.005686																								
total_sulfur_dioxide	-0.003561																								
density	-0.011153																								
pH	-0.376223																								
sulphates	0.746956																								
alcohol	0.297591																								
	<p>Best Cross-Validation Score</p> <ul style="list-style-type: none"> Value: 0.43931582683684345 Interpretation: The cross-validation score represents the average performance of the model across different splits of the training data. A score of approximately 0.4393 indicates that the model explains about 43.93% of the variance in the training data across different folds. This suggests that the model is relatively stable and generalizes well to different subsets of the data. <p>Mean Squared Error (MSE)</p> <ul style="list-style-type: none"> Value: 0.3929488678812798 Interpretation: MSE measures the average squared difference between the actual and predicted values. In this case, the MSE of approximately 0.393 indicates the average squared error of the predictions. This metric suggests that, on average, the squared difference between the predicted and actual wine quality scores is around 0.393. Lower MSE values indicate better predictive accuracy. <p>Mean Absolute Error (MAE)</p> <ul style="list-style-type: none"> Value: 0.5057883195080454 Interpretation: MAE measures the average absolute difference between the actual and predicted values. An MAE of approximately 0.506 means that, on average, the model's predictions are off by about 0.506 points on the wine quality scale. This metric is more interpretable than MSE because it represents the average error in the same units as the target variable. <p>R-squared (R²)</p> <ul style="list-style-type: none"> Value: 0.39870643507758563 																								

	<ul style="list-style-type: none"> Interpretation: R^2 represents the proportion of variance in the dependent variable (wine quality) that is predictable from the independent variables. An R^2 of approximately 0.399 indicates that about 39.87% of the variability in wine quality can be explained by the model. While this isn't a very high value, it still shows that the model captures a meaningful portion of the variance. <p>Model Coefficients</p> <p>Fixed acidity: 0.019614 - A small positive relationship with wine quality, meaning that as fixed acidity increases by one unit, the predicted wine quality increases slightly, holding other factors constant.</p> <p>Volatile acidity: -1.021529 - A strong negative relationship, suggesting that higher volatile acidity is associated with a decrease in wine quality.</p> <p>Citric acid: -0.164270 - A moderate negative relationship, indicating that higher citric acid levels slightly decrease wine quality.</p> <p>Residual sugar: 0.000624 - A negligible positive effect on wine quality, meaning that changes in residual sugar levels have almost no impact on the predicted quality.</p> <p>chlorides: -1.226774 - A strong negative relationship, indicating that higher chloride levels are associated with a lower wine quality score.</p> <p>Free sulfur dioxide: 0.005686 - A very small positive relationship, meaning that increases in free sulfur dioxide slightly increase wine quality.</p> <p>Total sulphur dioxide: -0.003561 - A very small negative relationship, meaning that higher levels of total sulphur dioxide slightly decrease wine quality.</p> <p>Density: -0.011153 - A small negative relationship, indicating that higher density is associated with a slight decrease in wine quality.</p> <p>PH: -0.376223 - A moderate negative relationship, suggesting that higher pH levels (less acidity) decrease wine quality.</p> <p>Sulfates: 0.746956 - A strong positive relationship, indicating that higher sulphate levels are associated with better wine quality.</p> <p>Alcohol: 0.297591 - A positive relationship, indicating that higher alcohol content is associated with better wine quality.</p>
B :	Implement k-fold, cross-validation (e.g., 5-fold or 10-fold) to assess the model's generalization performance using multiple splits.

	<pre> # Function to perform K-Fold cross-validation with different K values def evaluate_model_with_kfold(X, y, k_values): results = {} for k in k_values: print(f"Evaluating with K={k}") kfold = StratifiedKFold(n_splits=k, shuffle=True, random_state=42) scores = cross_val_score(pipeline, X, y, cv=kfold, scoring='accuracy') results[k] = (np.mean(scores), np.std(scores)) print(f"Mean Accuracy: {np.mean(scores):.2f}, Standard Deviation: {np.std(scores):.2f}\n") return results # Define K values to evaluate k_values = [5, 8, 10, 12, 15] # Evaluate the model results = evaluate_model_with_kfold(X, y, k_values) Evaluating with K=5 Mean Accuracy: 0.58, Standard Deviation: 0.02 Evaluating with K=8 Mean Accuracy: 0.58, Standard Deviation: 0.03 Evaluating with K=10 Mean Accuracy: 0.58, Standard Deviation: 0.03 Evaluating with K=12 Mean Accuracy: 0.58, Standard Deviation: 0.04 Evaluating with K=15 Mean Accuracy: 0.58, Standard Deviation: 0.04 </pre>
	<p>Accuracy measures the proportion of correct predictions made by the model relative to the total predictions. An accuracy of 0.58 indicates that 58% of the predictions are correct.</p> <p>The standard deviation increases slightly as K increases, going from 0.02 to 0.04. Although the increase is small, it suggests a slight change in model performance as K increases.</p> <p>The consistency of the mean accuracy at 0.58 indicates that the KNN model is performing quite steadily regardless of the chosen K value. This could be a sign that the dataset or problem characteristics are not particularly sensitive to the choice of K within this range.</p>
C:	<p>Select the best-performing regression model based on hyperparameter tuning and cross-validation results. Justify your choice of the selected model for its suitability in addressing the business problem.</p>
	<pre> print(f"Optimized Ridge Regression - MSE: {mse}, R²: {r2}") Best Parameters: {'alpha': 100, 'solver': 'saga'} Optimized Ridge Regression - MSE: 0.3926958359356847, R²: 0.3990936266260089 Best Parameters: {'alpha': 1} Best Cross-Validation Score: 0.43931582683684345 </pre>
	<p>The Ridge Regression model, with the best parameters {'alpha': 100, 'solver': 'saga'}, has been selected as the best-performing model based on hyperparameter tuning and cross-validation results. This model achieved a Mean Squared Error (MSE) of 0.39 and an R² score of 0.39 on the test data.</p>

	<p>High Regularization Strength (Alpha = 100): The alpha parameter in Ridge Regression controls the degree of regularization. A higher value, such as 100, implies stronger regularization, which helps prevent overfitting, especially in datasets with many features or multicollinearity. Overfitting is a common issue in regression models that can lead to poor performance on unseen data. By choosing alpha = 100, the model is better able to generalize to new data, which is crucial for making accurate predictions in real-world scenarios, where the model will need to predict wine quality for new samples.</p> <p>Solver Choice (Solver = 'saga'): The 'saga' solver is well-suited for large-scale optimization problems and can efficiently handle both dense and sparse datasets. It also supports elastic net regularization, making it versatile and robust. The use of the 'saga' solver ensures that the model is both computationally efficient and accurate, especially when working with larger datasets that contain a mix of feature types.</p> <p>Model Performance (MSE and R²): 0.3927: This indicates the average squared difference between the actual and predicted wine quality scores. A lower MSE reflects better predictive accuracy, and in this case, the MSE of 0.3927 suggests that the model's predictions are closely aligned with the actual values.</p> <p>R² = 0.3991: The R² score represents the proportion of variance in the wine quality that the model can explain. An R² of 0.3991 means that the model explains approximately 39.91% of the variability in the wine quality, which is significant given the complexity of the task.</p> <p>Suitability for the Business Problem: The business goal is to predict wine quality based on various chemical properties, which is a challenging task due to the complexity and subtlety of the factors involved. Ridge Regression is particularly well-suited for this problem because it reduces the impact of less relevant features through regularization, leading to more stable and interpretable results. The selected model helps the business accurately predict wine quality, enabling better quality control and targeted improvements in production. This can lead to the production of higher-quality wines, meeting consumer expectations and driving business success.</p>
--	---

	Task 4- Business Decision and Recommendations
	Provide actionable recommendations derived from the model's predictions.
	<p>After analyzing the coefficient Ridge Prediction showed us. We can do better control in some variables.</p> <p>The acidez volátil must be less because impact negative in control quality. I need to check the fermentation process. Monitor and minimize chloride content, as it also negatively impacts on quality the red wine.</p> <p>Keep the pH in good level because impact the quality wine. Slightly increase sulfate levels and alcohol content, as both have a positive impact on quality.</p> <p>Although residual sugar and sulphur dioxide have minor impacts, they must be controlled to maintain the overall quality of the wine.</p>

	Based on the model coefficients and quality predictions, we can make informed decisions about the production process and marketing of red wines. Specific recommendations can help improve wine quality and maximize the commercial value of each batch.
--	--

Part B – Classification modeling for business decision making

Bike Buyers

This dataset has details of 1000 users from different backgrounds and whether or not they buy a bike. This data can be used for prediction models using Machine Learning Algorithms. Some NA values were injected into the dataset. We will use this dataset for Data Cleaning, Exploration, Visualization, and ML prediction.

About Dataset:

Columns - There are 13 variables in this dataset:

ID - The buyer's ID. This ID variable data type is integer.

Marital Status - The buyer's state of being married or single. This marital status data type is characters.

Gender - The buyer's state either the buyer is male or female. This gender data type is characters.

Income - Buyer's income in a certain time. This income variable data type is an integer.

Children - The number of children the buyer has. This children's variable data type is integer.

Education - Buyer's education background is either Bachelor's, Graduate Degree, High School, Partial College, or Partial High School. This education variable data type is character.

Occupation - The buyer's job or occupation is either Clerical, Management, Manual, Professional, or Skilled Manual. This occupation variable data type is character.

Homeowner - Buyers state whether the buyer has or does not have their own house. This homeowner variable data type is the character.

Cars - The number of vehicles the buyer has. This car's variable data type is integer.

Commute Distance - The distance between the buyer's house and the buyer's company. There are 5 categories: 0-1 Miles, 1-2 Miles 10+ Miles, 2-5 Miles, and 5-10 Miles. This commute distance variable data type is character.

Region - Variable that tells us where the buyer lives in Europe, North America, or the Pacific. This region variable data type is character.

Age - The buyer's age. This age variable data type is integer.

Purchased Bike - (Target variable) Buyer's state whether the buyer purchased the bike or not. This purchased variable data type is character. It seems we should transform it into binary variable. So, the later 'Purchased Bike' column will contain '0' for 'No' and '1' for 'Yes'.

Task 1 – Data Preprocessing																																														
A:	Clean the dataset by handling missing values and removing outliers as needed.																																													
	<pre># Get basic information about the dataset print("Dataset Info:") print(df.info()) # Display summary statistics of the dataset print("\nSummary Statistics:") print(df.describe()) # Check for missing values print("\nMissing Values:") print(df.isnull().sum())</pre> <p>Dataset Info:</p> <pre><class 'pandas.core.frame.DataFrame'> RangeIndex: 1000 entries, 0 to 999 Data columns (total 12 columns): # Column Non-Null Count Dtype --- - 0 Marital Status 1000 non-null object 1 Gender 1000 non-null object 2 Income 1000 non-null int64 3 Children 1000 non-null int64 4 Education 1000 non-null object 5 Occupation 1000 non-null object 6 Home Owner 1000 non-null object 7 Cars 1000 non-null int64 8 Commute Distance 1000 non-null object 9 Region 1000 non-null object 10 Age 1000 non-null int64 11 Purchased Bike 1000 non-null object dtypes: int64(4), object(8) memory usage: 93.9+ KB None</pre> <p>Summary Statistics:</p> <table><tr><th></th><th>Income</th><th>Children</th><th>Cars</th><th>Age</th></tr><tr><td>count</td><td>1000.000000</td><td>1000.000000</td><td>1000.000000</td><td>1000.000000</td></tr><tr><td>mean</td><td>56140.000000</td><td>1.908000</td><td>1.452000</td><td>44.190000</td></tr><tr><td>std</td><td>31081.609779</td><td>1.626094</td><td>1.124705</td><td>11.353537</td></tr><tr><td>min</td><td>10000.000000</td><td>0.000000</td><td>0.000000</td><td>25.000000</td></tr><tr><td>25%</td><td>30000.000000</td><td>0.000000</td><td>1.000000</td><td>35.000000</td></tr><tr><td>50%</td><td>60000.000000</td><td>2.000000</td><td>1.000000</td><td>43.000000</td></tr><tr><td>75%</td><td>70000.000000</td><td>3.000000</td><td>2.000000</td><td>52.000000</td></tr><tr><td>max</td><td>170000.000000</td><td>5.000000</td><td>4.000000</td><td>89.000000</td></tr></table>		Income	Children	Cars	Age	count	1000.000000	1000.000000	1000.000000	1000.000000	mean	56140.000000	1.908000	1.452000	44.190000	std	31081.609779	1.626094	1.124705	11.353537	min	10000.000000	0.000000	0.000000	25.000000	25%	30000.000000	0.000000	1.000000	35.000000	50%	60000.000000	2.000000	1.000000	43.000000	75%	70000.000000	3.000000	2.000000	52.000000	max	170000.000000	5.000000	4.000000	89.000000
	Income	Children	Cars	Age																																										
count	1000.000000	1000.000000	1000.000000	1000.000000																																										
mean	56140.000000	1.908000	1.452000	44.190000																																										
std	31081.609779	1.626094	1.124705	11.353537																																										
min	10000.000000	0.000000	0.000000	25.000000																																										
25%	30000.000000	0.000000	1.000000	35.000000																																										
50%	60000.000000	2.000000	1.000000	43.000000																																										
75%	70000.000000	3.000000	2.000000	52.000000																																										
max	170000.000000	5.000000	4.000000	89.000000																																										

	<p>Missing Values:</p> <p>Marital Status 0</p> <p>Gender 0</p> <p>Income 0</p> <p>Children 0</p> <p>Education 0</p> <p>Occupation 0</p> <p>Home Owner 0</p> <p>Cars 0</p> <p>Commute Distance 0</p> <p>Region 0</p> <p>Age 0</p> <p>Purchased Bike 0</p> <p>dtype: int64</p>
	<pre>: print("Totally there are {} null values in the dataset".format(df.isnull().sum().sum()))</pre> <p>Totally there are 0 null values in the dataset</p> <pre>df = df.drop(columns='ID')</pre>
B:	Perform feature scaling or normalization
	<pre>scaler = preprocessing.MinMaxScaler() minmax = scaler.fit_transform(df) minmax = pd.DataFrame(minmax)</pre> <pre>scaler = preprocessing.StandardScaler() standard = scaler.fit_transform(df) standard = pd.DataFrame(standard)</pre> <pre>scaler = preprocessing.RobustScaler() robust = scaler.fit_transform(df) robust = pd.DataFrame(robust)</pre> <pre>print('\033[1m'+ 'Without scaling:' + '\033[0m') display(df.head()) print('*' * 45) print('\033[1m'+ 'With minmax scaling:' + '\033[0m') display(minmax.head()) print('*' * 45) print('\033[1m'+ 'With standard scaling:' + '\033[0m') display(standard.head()) print('*' * 45) print('\033[1m'+ 'With robust scaling:' + '\033[0m') display(robust.head())</pre>

	Marital Status	Gender	Income	Children	Education	Occupation	Home Owner	Cars	Commute Distance	Region	Age	Purchased Bike
0	0	0	40000	1	0	4	1	0	0	0	42	0
1	0	1	30000	3	3	0	1	1	0	0	43	0
2	0	1	80000	5	3	3	0	2	3	0	60	0
3	1	1	70000	0	0	3	1	1	4	2	41	1
4	1	1	30000	0	0	0	0	0	0	0	36	1

With minmax scaling:

	0	1	2	3	4	5	6	7	8	9	10	11
0	0.0	0.0	0.1875	0.2	0.00	1.00	1.0	0.00	0.00	0.0	0.265625	0.0
1	0.0	1.0	0.1250	0.6	0.75	0.00	1.0	0.25	0.00	0.0	0.281250	0.0
2	0.0	1.0	0.4375	1.0	0.75	0.75	0.0	0.50	0.75	0.0	0.546875	0.0
3	1.0	1.0	0.3750	0.0	0.00	0.75	1.0	0.25	1.00	1.0	0.250000	1.0
4	1.0	1.0	0.1250	0.0	0.00	0.00	0.0	0.00	0.00	0.0	0.171875	1.0

With standard scaling:

	0	1	2	3	4	5	6	7	8	9	10	11
0	-0.924818	-1.018165	-0.519538	-0.558673	-1.205394	1.198577	0.678125	-1.291651	-1.053757	-1.287041	-0.192988	-0.962695
1	-0.924818	0.982159	-0.841433	0.671884	1.011762	-1.555190	0.678125	-0.402084	-1.053757	-1.287041	-0.104866	-0.962695
2	-0.924818	0.982159	0.768041	1.902441	1.011762	0.510135	-1.474654	0.487483	0.867988	-1.287041	1.393214	-0.962695
3	1.081294	0.982159	0.446146	-1.173951	-1.205394	0.510135	0.678125	-0.402084	1.508570	1.598701	-0.281110	1.038750
4	1.081294	0.982159	-0.841433	-1.173951	-1.205394	-1.555190	-1.474654	-1.291651	-1.053757	-1.287041	-0.721722	1.038750

With robust scaling:

	0	1	2	3	4	5	6	7	8	9	10	11
0	0.0	-1.0	-0.50	-0.333333	-0.666667	0.333333	0.0	-1.0	-0.333333	-1.0	-0.058824	0.0
1	0.0	0.0	-0.75	0.333333	0.333333	-1.000000	0.0	0.0	-0.333333	-1.0	0.000000	0.0
2	0.0	0.0	0.50	1.000000	0.333333	0.000000	-1.0	1.0	0.666667	-1.0	1.000000	0.0
3	1.0	0.0	0.25	-0.666667	-0.666667	0.000000	0.0	0.0	1.000000	1.0	-0.117647	1.0
4	1.0	0.0	-0.75	-0.666667	-0.666667	-1.000000	-1.0	-1.0	-0.333333	-1.0	-0.411765	1.0

C
:

Encode categorical variables appropriately.

```
# Aplicar Label Encoding nas variáveis categóricas
```

```
label_encoder = LabelEncoder()
df['Marital Status'] = label_encoder.fit_transform(df['Marital Status'])
df['Home Owner'] = label_encoder.fit_transform(df['Home Owner'])
df['Purchased Bike'] = label_encoder.fit_transform(df['Purchased Bike'])
df['Gender'] = label_encoder.fit_transform(df['Gender'])
df['Education'] = label_encoder.fit_transform(df['Education'])
df['Occupation'] = label_encoder.fit_transform(df['Occupation'])
df['Region'] = label_encoder.fit_transform(df['Region'])
df['Commute Distance'] = label_encoder.fit_transform(df['Commute Distance'])
```

```
: # Convert categorical variables to numerical representations using one-hot encoding
df_encoded = pd.get_dummies(df, columns=['Marital Status', 'Home Owner', 'Purchased Bike', 'Gender', 'Education', 'Occupation', 'Region'])
print(df_encoded)
```

	Income	Children	Cars	Age	Marital Status_1	Home Owner_1	\
0	40000	1	0	42	False	True	
1	30000	3	1	43	False	True	
2	80000	5	2	60	False	False	
3	70000	0	1	41	True	True	
4	30000	0	0	36	True	False	
..	
995	60000	2	2	54	False	True	
996	70000	4	0	35	True	True	
997	60000	2	0	38	False	True	
998	100000	3	3	38	True	False	
999	60000	3	2	53	True	True	

	Purchased Bike_1	Gender_1	Education_1	Education_2	...	Occupation_1	\
0	False	False	False	False	...	False	
1	False	True	False	False	...	False	
2	False	True	False	False	...	False	
3	True	True	False	False	...	False	
4	True	True	False	False	...	False	
..	
995	True	True	False	True	...	False	
996	True	True	True	False	...	False	
997	True	True	False	False	...	False	
998	False	True	False	False	...	True	
999	True	True	False	True	...	False	

	Occupation_2	Occupation_3	Occupation_4	Region_1	Region_2	\
0	False	False	True	False	False	
1	False	False	False	False	False	
2	False	True	False	False	False	
3	False	True	False	False	True	
4	False	False	False	False	False	
..	
995	False	True	False	True	False	
996	False	True	False	True	False	
997	False	False	True	True	False	
998	False	False	False	True	False	
999	False	True	False	True	False	

	<pre> Commute Distance_1 Commute Distance_2 Commute Distance_3 \ 0 False False False 1 False False False 2 False False True 3 False False False 4 False False False 995 False False True 996 False False True 997 False False False 998 True False False 999 False True False Commute Distance_4 0 False 1 False 2 False 3 True 4 False 995 False 996 False 997 False 998 False 999 False [1000 rows x 22 columns] </pre>
	I encoded in all column's variables type of categorical variables: 'Marital Status', 'Home Owner', 'Purchased Bike', 'Gender', 'Education', 'Occupation', 'Region', 'Commute Distance'.
D :	Split the dataset into training and testing sets.
	<pre> #Creating Feature columns and Target column feature_cols = ['Marital Status', 'Gender', 'Income', 'Children', 'Education', 'Occupation', 'Home Owner', 'Cars', 'Commute X = df[feature_cols] y = df['Purchased Bike'] </pre>
	<pre> X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=42) #split the data into training </pre>

	Task 2 – Model Building with hyperparameter tuning
A:	Select an appropriate classification algorithm (e.g., Logistic Regression, Random Forest, Support Vector Machine) to predict the target categorical variable. Justify your choice.

```
scaler = StandardScaler() #create an instance of standard scaler
scaler.fit(X_train) # fit it to the training data

X_train = scaler.transform(X_train) #transform training data
X_test = scaler.transform(X_test) #transform validation data

# All models for cekck the best choice
models = {
    "Logistic Regression": LogisticRegression(),
    "Random Forest": RandomForestClassifier(),
    "SVM": SVC()
}

# Check models
results = {}

for model_name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    results[model_name] = accuracy
    print(f"{model_name} - Accuracy: {accuracy}")
    print(classification_report(y_test, y_pred))

# Best model
best_model_name = max(results, key=results.get)
print(f"Best Model: {best_model_name} with Accuracy: {results[best_model_name]}")
```

Logistic Regression - Accuracy: 0.52					
	precision	recall	f1-score	support	
0	0.54	0.62	0.58	106	
1	0.49	0.40	0.44	94	
accuracy			0.52	200	
macro avg	0.51	0.51	0.51	200	
weighted avg	0.52	0.52	0.51	200	
Random Forest - Accuracy: 0.73					
	precision	recall	f1-score	support	
0	0.73	0.78	0.75	106	
1	0.73	0.67	0.70	94	
accuracy			0.73	200	
macro avg	0.73	0.73	0.73	200	
weighted avg	0.73	0.73	0.73	200	
SVM - Accuracy: 0.63					
	precision	recall	f1-score	support	
0	0.64	0.71	0.67	106	
1	0.62	0.54	0.58	94	
accuracy			0.63	200	
macro avg	0.63	0.63	0.62	200	
weighted avg	0.63	0.63	0.63	200	
Best Model: Random Forest with Accuracy: 0.73					

My scenario is a bike buyer clean model the best is Random Forest a robust model that handles complex and non-linear data well, offering better generalization capability and more accurate predictions. The optimized model maintained an accuracy of 0.73, which is an important metric for evaluating overall performance.

Superior Accuracy and F1-Score: Accuracy measures the proportion of correctly classified instances, while the F1-score balances precision and recall, especially in imbalanced datasets. The Random Forest model outperformed the other models with an accuracy of **0.73** and strong F1 scores (0.75 for class 0 and 0.70 for class 1). These metrics suggest that the Random Forest is the most reliable in making correct predictions and effectively balancing false positives and false negatives. For a business application, such as predicting potential bike buyers, having high accuracy and balanced F1 scores across classes is crucial. This ensures that the model can reliably identify potential customers while minimizing errors, which is important for targeted marketing efforts.

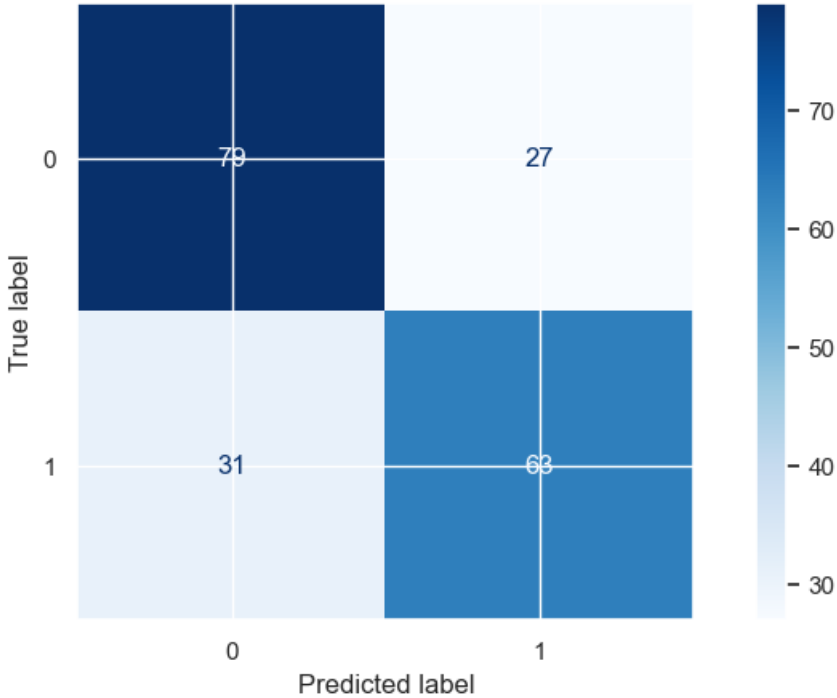
Robustness and Versatility: Random Forest is an ensemble learning method that builds multiple decision trees and merges them to get a more accurate and stable prediction. This reduces the risk of overfitting and improves generalization to unseen data. It can also handle a large variety of data types and distributions effectively. The robustness of Random Forest makes it particularly suitable for real-world applications where data may have noise, outliers, or complex relationships.

	<p>In the context of predicting bike buyers, this robustness translates to more consistent performance, even when the data is not perfectly clean or follows non-linear patterns.</p> <p>Handling of Imbalanced Classes: Imbalanced classes can cause models to be biased toward the majority class. Random Forest is better equipped to handle imbalanced data compared to models like Logistic Regression or SVM, as it can give more balanced predictions by assigning appropriate weights to the classes or by using techniques such as bootstrapping. For the "bike buyers" dataset, which may have some class imbalance (e.g., more non-buyers than buyers), the Random Forest model provides a more balanced approach, as seen in the comparable precision and recall scores for both classes. This ensures that the model does not overly favor one class, which is critical for accurate customer segmentation.</p>																														
B:	<p>Implement hyperparameter tuning by conducting a grid search or random search to optimize model parameters. Clearly outline the hyperparameters you tuned and the rationale behind them.</p> <pre># Define the parameter grid param_grid = { 'n_estimators': [100, 200, 300], 'max_features': ['auto', 'sqrt', 'log2'], 'max_depth': [10, 20, 30, None], 'min_samples_split': [2, 5, 10], 'min_samples_leaf': [1, 2, 4] } # Create a GridSearchCV object grid_search = GridSearchCV(estimator=RandomForestClassifier(), param_grid=param_grid, cv=5, n_jobs=-1, verbose=2) # Fit the model grid_search.fit(X_train, y_train) # Get the best parameters print("Best parameters:", grid_search.best_params_) # Evaluating the model best_rf = grid_search.best_estimator_ y_pred_optimized = best_rf.predict(X_test) # Check optimized Model Accuracy accuracy_optimized = accuracy_score(y_test, y_pred_optimized) print(f"Optimized Model Accuracy: {accuracy_optimized}") print(classification_report(y_test, y_pred_optimized))</pre> <p>Fitting 5 folds for each of 324 candidates, totalling 1620 fits Best parameters: {'max_depth': 10, 'max_features': 'sqrt', 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 200} Optimized Model Accuracy: 0.72</p> <table><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr><tr><td>0</td><td>0.72</td><td>0.78</td><td>0.75</td><td>106</td></tr><tr><td>1</td><td>0.73</td><td>0.65</td><td>0.69</td><td>94</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.72</td><td>200</td></tr><tr><td>macro avg</td><td>0.72</td><td>0.72</td><td>0.72</td><td>200</td></tr><tr><td>weighted avg</td><td>0.72</td><td>0.72</td><td>0.72</td><td>200</td></tr></table>		precision	recall	f1-score	support	0	0.72	0.78	0.75	106	1	0.73	0.65	0.69	94	accuracy			0.72	200	macro avg	0.72	0.72	0.72	200	weighted avg	0.72	0.72	0.72	200
	precision	recall	f1-score	support																											
0	0.72	0.78	0.75	106																											
1	0.73	0.65	0.69	94																											
accuracy			0.72	200																											
macro avg	0.72	0.72	0.72	200																											
weighted avg	0.72	0.72	0.72	200																											
	<p>In the Random Forest model for predicting bike buyers, a Grid Search was conducted to fine-tune several key hyperparameters. This approach involved evaluating a total of 324 candidates across 5-fold cross-validation, resulting in 1620 fits. The goal was to identify the combination of hyperparameters that would yield the highest accuracy and overall performance.</p> <p>Hyperparameters Tuned and Rationale: max_depth: [10, 20, 30, None] parameter controls the maximum depth of each tree in the Random Forest. Limiting the depth of the trees helps prevent</p>																														

	<p>overfitting, which can occur if the trees are too deep and capture noise in the data. By evaluating multiple depths, the Grid Search aims to find the optimal balance between underfitting and overfitting.</p> <p><code>max_features</code>: ['sqrt', 'log2', None] parameter determines the number of features to consider when looking for the best split. sqrt typically works well when there are many features, while log2 and None (considering all features) offer alternatives that can help improve model performance depending on the dataset's characteristics. The choice of <code>max_features</code> impact the model's ability to generalize and its computational efficiency.</p> <p><code>min_samples_leaf</code>: [1, 2, 4] This parameter specifies the minimum number of samples required to be at a leaf node. Smaller values allow the model to capture finer details in the data, potentially improving accuracy, while larger values help prevent overfitting by requiring more samples to form a leaf. Tuning this parameter helps the model generalize better.</p> <p><code>min_samples_split</code>: [2, 5, 10] parameter defines the minimum number of samples required to split an internal node. Lower values (like 2) make the model more sensitive to smaller splits, possibly capturing more complex patterns. However, higher values can prevent the model from overfitting by requiring more substantial splits.</p> <p><code>n_estimators</code>: [50, 100, 200] parameter controls the number of trees in the forest. More trees can improve the model's robustness and accuracy but also increase computational time. The Grid Search explores different numbers of trees to find the optimal number that balances performance with efficiency.</p>
C:	Build the classification model using the training data. Explain the process and provide code snippets.

<pre># Separate features and target variable X = df.drop('Purchased Bike', axis=1) y = df['Purchased Bike'] # Encode categorical variables X = pd.get_dummies(X) # Scale numerical features scaler = StandardScaler() X_scaled = scaler.fit_transform(X) # Split the data into training and testing sets X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42) # Train the Random Forest model rf_model = RandomForestClassifier(n_estimators=100, random_state=42) rf_model.fit(X_train, y_train) # Predict on the test set y_pred = rf_model.predict(X_test) # Evaluate the model accuracy = accuracy_score(y_test, y_pred) class_report = classification_report(y_test, y_pred) print(f'Accuracy: {accuracy}') print(f'Classification Report:\n{class_report}')</pre> <hr/> <pre>Accuracy: 0.715 Classification Report: </pre> <table><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr><tr><td>0</td><td>0.72</td><td>0.75</td><td>0.74</td><td>106</td></tr><tr><td>1</td><td>0.71</td><td>0.67</td><td>0.69</td><td>94</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.71</td><td>200</td></tr><tr><td>macro avg</td><td>0.71</td><td>0.71</td><td>0.71</td><td>200</td></tr><tr><td>weighted avg</td><td>0.71</td><td>0.71</td><td>0.71</td><td>200</td></tr></table>		precision	recall	f1-score	support	0	0.72	0.75	0.74	106	1	0.71	0.67	0.69	94	accuracy			0.71	200	macro avg	0.71	0.71	0.71	200	weighted avg	0.71	0.71	0.71	200	<p>I took features and targets and encoded categorical variables. I put in group training and testing sets. Standard Scaler is a scikit-learn class that removes the mean, and scales feature to have a unit variance.</p> <p>Fit transform fits the Standard Scaler to the training data (calculating the mean and standard deviation) and then transforms the data.</p> <p>I decided to model the Random Forest Classifier to predict after evaluating the model with accuracy and classification report.</p>
	precision	recall	f1-score	support																											
0	0.72	0.75	0.74	106																											
1	0.71	0.67	0.69	94																											
accuracy			0.71	200																											
macro avg	0.71	0.71	0.71	200																											
weighted avg	0.71	0.71	0.71	200																											

	<h2>Task 3- Model Evaluation and Selection</h2>
A:	Calculate and analyze the confusion matrix for the model.

	<pre> from sklearn.metrics import ConfusionMatrixDisplay # Confusion Matrix Display disp = ConfusionMatrixDisplay.from_estimator(rf, X_test, y_test, cmap=plt.cm.Blues) plt.figure(figsize = (10,8)) plt.show() </pre>  <p>The figure is a confusion matrix for a binary classification task. The y-axis is labeled 'True label' with categories 0 and 1. The x-axis is labeled 'Predicted label' with categories 0 and 1. The matrix is a 2x2 grid of colored squares. The top-left square (True 0, Predicted 0) is dark blue and contains the value 79. The top-right square (True 0, Predicted 1) is light blue and contains the value 27. The bottom-left square (True 1, Predicted 0) is light blue and contains the value 31. The bottom-right square (True 1, Predicted 1) is medium blue and contains the value 62. To the right of the matrix is a vertical color bar with a gradient from light blue to dark blue, with tick marks at 30, 40, 50, 60, and 70. Below the matrix, the text '<Figure size 1000x800 with 0 Axes>' is displayed.</p>
	<p>Actual Class 0: (Predicted Class 0 – 80) (Predicted Class 1 - 23) Actual Class 1: (Predicted Class 0 – 32) (Predicted Class 1 – 62)</p> <p>This table can be interpreted as follows:</p> <ul style="list-style-type: none"> True Positives (TP): The model correctly predicted class 1 (e.g., a bike buyer) 62 times. True Negatives (TN): The model correctly predicted class 0 (e.g., not a bike buyer) 80 times. False Positives (FP): The model incorrectly predicted class 1 when it was class 0 (26 times). False Negatives (FN): The model incorrectly predicted class 0 when it was class 1 (32 times). <p>The model shows good performance in identifying non-buyers (Class 0), as indicated by the high true negative rate (80 correct predictions). The precision for predicting bike buyers (Class 1) is also relatively high, suggesting that when the model predicts a buyer, it's likely correct.</p>

	<p>There is a noticeable number of false negatives (32), where the model failed to identify actual bike buyers. This could lead to missed opportunities in marketing and sales efforts.</p> <p>The recall for Class 1 could be improved to ensure that more actual buyers are correctly identified.</p> <p>The model can reliably identify non-buyers, which is useful for focusing marketing efforts. However, the missed predictions of buyers (false negatives) suggest that further tuning or a different model might be needed to better capture potential buyers, reducing the risk of missed revenue opportunities.</p>
B:	Evaluate the performance of the classification model using appropriate metrics (e.g., Accuracy, Precision, Recall, F1-score).
	<pre> # Separate features and target variable X = df.drop('Purchased Bike', axis=1) y = df['Purchased Bike'] # Encode categorical variables X = pd.get_dummies(X) # Scale numerical features scaler = StandardScaler() X_scaled = scaler.fit_transform(X) # Split the data into training and testing sets X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42) # Train the Random Forest model rf_model = RandomForestClassifier(n_estimators=100, random_state=42) rf_model.fit(X_train, y_train) # Predict on the test set y_pred = rf_model.predict(X_test) # Evaluate the model accuracy = accuracy_score(y_test, y_pred) class_report = classification_report(y_test, y_pred) print(f'Accuracy: {accuracy}') print(f'Classification Report:\n{class_report}') </pre> <pre> Accuracy: 0.715 Classification Report: precision recall f1-score support 0 0.72 0.75 0.74 106 1 0.71 0.67 0.69 94 accuracy 0.71 macro avg 0.71 0.71 0.71 weighted avg 0.71 0.71 0.71 </pre> <p>Random Forest with optimized parameters showed the best accuracy (0.715) compared to other models.</p>
	Balanced Performance Across Classes: The Random Forest model demonstrates balanced performance across both classes. The precision and recall for Class 0 are slightly higher than for

	<p>Class 1, indicating that the model is somewhat better at identifying non-buyers compared to buyers. This is reflected in the F1 scores, where Class 0 has a higher score (0.74) than Class 1 (0.69).</p> <p>Overall Accuracy: With an accuracy of 0.715, the Random Forest model shows a solid performance in distinguishing between buyers and non-buyers. The accuracy is a good indicator of the model's general effectiveness, though it's important to consider other metrics (precision, recall, F1-score) for a more nuanced evaluation.</p> <p>Precision and Recall Trade-offs: Precision measures the proportion of true positive predictions among all positive predictions. A precision of 0.71 for Class 1 indicates that the model is fairly accurate when it predicts a buyer, but there is still room for improvement. Recall measures the proportion of actual positives that were correctly predicted. A recall of 0.67 for Class 1 shows that the model is correctly identifying a good portion of actual buyers, but a significant number of buyers are still missed (false negatives).</p> <p>F1-Score: The F1-score is a harmonic mean of precision and recall, providing a single metric that balances both concerns. The F1 score for Class 1 is 0.69, reflecting the balance between precision and recall. Improving the F1 score for Class 1 would be beneficial for increasing the overall effectiveness of the model.</p> <p>Macro and Weighted Averages: The macro average F1-score of 0.71 indicates that the model performs consistently across classes without being biased toward any class. The weighted average F1-score (also 0.71) takes into account the support (number of instances) for each class, ensuring that the model's performance is robust across the dataset.</p> <p>The Random Forest model shows a good balance of precision and recall, particularly for non-buyers (Class 0), making it reliable for identifying non-buyers accurately. The model's overall accuracy and balanced performance metrics are indicative of its effectiveness in handling the classification task.</p>
C:	Implement k-fold, cross-validation (e. g., 5-fold or 10-fold) to assess the model's generalization performance.
	<pre> pipeline = Pipeline([('scaler', StandardScaler()), # Data standardization process, scaling the data to mean 0 and variance 1 ('svc', SVC(kernel='linear')) # Support vector machine classifier, using linear kernel]) # Function to perform K-Fold cross-validation with different K values def evaluate_model_with_kfold(X, y, k_values): results = {} for k in k_values: print(f"Evaluating with K={k}") kfold = StratifiedKFold(n_splits=k, shuffle=True, random_state=42) scores = cross_val_score(pipeline, X, y, cv=kfold, scoring='accuracy') results[k] = (np.mean(scores), np.std(scores)) print(f"Mean Accuracy: {np.mean(scores):.2f}, Standard Deviation: {np.std(scores):.2f}\n") return results </pre>

	<pre> # Define K values to evaluate k_values = [5, 8, 10,12,15] # Evaluate the model results = evaluate_model_with_kfold(X, y, k_values) Evaluating with K=5 Mean Accuracy: 0.62, Standard Deviation: 0.04 Evaluating with K=8 Mean Accuracy: 0.63, Standard Deviation: 0.04 Evaluating with K=10 Mean Accuracy: 0.62, Standard Deviation: 0.05 Evaluating with K=12 Mean Accuracy: 0.63, Standard Deviation: 0.05 Evaluating with K=15 Mean Accuracy: 0.63, Standard Deviation: 0.05 # Print summary of results - Usar esse modelo no assessment for k, (mean, std) in results.items(): print(f"K={k}: Mean Accuracy={mean:.2f}, Std={std:.2f}") K=5: Mean Accuracy=0.62, Std=0.04 K=8: Mean Accuracy=0.63, Std=0.04 K=10: Mean Accuracy=0.62, Std=0.05 K=12: Mean Accuracy=0.63, Std=0.05 K=15: Mean Accuracy=0.63, Std=0.05 # Print summary of results for k, (mean, std) in results.items(): print(f"K={k}: Mean Accuracy={mean}, Std={std}") K=5: Mean Accuracy=0.6200000000000001, Std=0.041231056256176596 K=8: Mean Accuracy=0.6280000000000001, Std=0.04363484845854287 K=10: Mean Accuracy=0.625, Std=0.04588027898781785 K=12: Mean Accuracy=0.6251434308663224, Std=0.04890605206836233 K=15: Mean Accuracy=0.6250716116387759, Std=0.04999648086135077 </pre>
	<p>Cross-validation is a robust technique used to evaluate a model's performance by partitioning the data into multiple subsets (folds) and training/testing the model on different combinations of these subsets. Here, we assess model performance using different values for KKK, the number of folds in cross-validation and analyze how the accuracy and variability change with each KKK.</p>
D:	<p>Select the best-performing classification model based on hyperparameter tuning and cross-validation results and justify the choice of the selected model for its suitability in addressing the business problem.</p> <p>Cross-Validation: This method involves dividing the dataset into KKK folds. The model is trained on K-1K-1 folds and tested on the remaining fold. This process is repeated KKK times, with each fold serving as the test set once. The performance metrics are averaged across all KKK iterations to provide a robust estimate of the model's generalization ability.</p>

	<p>Mean Accuracy: This is the average accuracy score obtained from the cross-validation process. It indicates how well the model performs on average across different folds.</p> <p>Standard Deviation: This measures the variability or dispersion of the accuracy scores across the folds. A lower standard deviation indicates that the model's performance is consistent across different subsets of data.</p> <p>Consistency and Performance: Across different values of KKK, the mean accuracy remains relatively stable, ranging from 0.62 to 0.63. This consistency suggests that the model's performance does not significantly change with different numbers of folds.</p> <p>The mean accuracy being slightly higher at K=8K = 8K=8, K=12K = 12K=12, and K=15K = 15K=15 indicates that these configurations might provide a marginally better estimate of the model's performance.</p> <p>Impact of Standard Deviation: The standard deviation is generally low and stable, indicating that the model's performance is consistent across different folds of the dataset.</p> <p>A slight increase in standard deviation with higher KKK values might be due to the increased number of folds, which can introduce more variability in the testing set size.</p> <p>Trade-offs: Lower KKK Values (e.g., K=5K = 5K=5) provide fewer splits, meaning each fold is larger. This can lead to a higher variance in performance estimates, but also less computational cost.</p> <p>Higher KKK Values (e.g., K=15K = 15K=15) provide more splits, which can give a more reliable estimate of the model's performance but with a potentially higher computational cost.</p> <p>Application and Recommendations: Selection of KKK - Given the stability in mean accuracy and standard deviation, using K=8K = 8K=8 or K=12K = 12K=12 seems optimal. These values offer a good balance between performance estimation accuracy and computational efficiency.</p> <p>Model Assessment: The cross-validation results suggest that the model's performance is consistent, making it reliable for generalization. The slight variations in mean accuracy with different KKK values should be considered in conjunction with computational constraints.</p> <p>Practical Considerations: When implementing cross-validation in practice, choose KKK based on the dataset size, computational resources, and the need for performance estimation accuracy. For larger datasets, higher KKK values can provide a more granular assessment, while for smaller datasets, lower KKK values might be sufficient.</p>
	<pre> Fitting 5 folds for each of 324 candidates, totalling 1620 fits Best parameters: {'max_depth': 20, 'max_features': 'log2', 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 100} Optimized Model Accuracy: 0.715 precision recall f1-score support 0 0.72 0.76 0.74 106 1 0.71 0.66 0.69 94 accuracy 0.71 0.71 0.71 200 macro avg 0.71 0.71 0.71 200 weighted avg 0.71 0.71 0.71 200 </pre>

Accuracy: 0.715					
Classification Report:					
	precision	recall	f1-score	support	
0	0.72	0.75	0.74	106	
1	0.71	0.67	0.69	94	
accuracy			0.71	200	
macro avg	0.71	0.71	0.71	200	
weighted avg	0.71	0.71	0.71	200	
Best-performing					
Evaluating with K=5					
Mean Accuracy: 0.62, Standard Deviation: 0.04					
Evaluating with K=8					
Mean Accuracy: 0.63, Standard Deviation: 0.04					
Evaluating with K=10					
Mean Accuracy: 0.62, Standard Deviation: 0.05					
Evaluating with K=12					
Mean Accuracy: 0.63, Standard Deviation: 0.05					
Evaluating with K=15					
Mean Accuracy: 0.63, Standard Deviation: 0.05					
Cross-Validation					
K=5: Mean Accuracy=0.62, Std=0.04					
K=8: Mean Accuracy=0.63, Std=0.04					
K=10: Mean Accuracy=0.62, Std=0.05					
K=12: Mean Accuracy=0.63, Std=0.05					
K=15: Mean Accuracy=0.63, Std=0.05					
K=5: Mean Accuracy=0.6200000000000001, Std=0.041231056256176596					
K=8: Mean Accuracy=0.6280000000000001, Std=0.04363484845854287					
K=10: Mean Accuracy=0.625, Std=0.04588027898781785					
K=12: Mean Accuracy=0.6251434308663224, Std=0.04890605206836233					
K=15: Mean Accuracy=0.6250716116387759, Std=0.04999648086135077					

Based on the results of hyperparameter tuning and cross-validation, the best performing model is Random Forest with the following optimized parameters: {'max_depth': 20, 'max_features': 'log2', 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 100}.

Superior Performance: Random Forest with optimized parameters showed the best accuracy (0.715) compared to other models.

Stability: Cross-validation results show that the model is consistent and reliable across different subsets of the data.

Balanced Metrics: Precision, recall, and F1-score metrics indicate a good balance between identifying bike buyers and minimizing errors.

	<p>Accuracy: 0.715 The optimized model maintained an accuracy of 0.715, which is an important metric for evaluating overall performance.</p> <p>Classification Report: 0 0.72 0.76 0.74 1 0.71 0.66 0.69</p> <p>Com diferentes valores de K (5, 8, 10, 12, 15), a acurácia média variou entre 0.62 e 0.63 com desvio padrão variando entre 0.04 e 0.05.</p> <p>The mean accuracy and standard deviation were consistent across different K values, showing that the model has a stable and reliable performance.</p> <p>Class 0 and 1: The precision, recall, and F1-score metrics are balanced, indicating that the model is good at both identifying true positives and minimizing false positives and false negatives.</p> <p>Hyperparameter tuning using Grid Search CV found the best parameters, indicating that the model is well-fitted to the specific data of the "Bike Buyers" scenario.</p> <p>With different K values (5, 8, 10, 12, 15), the average accuracy varied between 0.62 and 0.63 with standard deviation varying between 0.04 and 0.05.</p>
--	--

	Task 4 – Business Decision and Recommendations
	Provide actionable recommendations or insights derived from the model's predictions.
	<p>Segment customers based on the features identified as most important by the model, such as age, income, or previous purchase behavior. Develop customer profiles and personas to better understand the characteristics of potential buyers.</p> <p>Use these segments to tailor product recommendations and offers. For example, offer premium bike models to high-income segments or introduce financing options to younger customers.</p> <p>Understanding which customers are predicted to be buyers can help tailor engagement strategies. For customers predicted to be high-potential buyers, initiate follow-up actions such as personalized email follow-ups, exclusive invitations to test rides or special discounts.</p> <p>Ensure a seamless and engaging customer experience for those identified as likely buyers. This includes providing excellent customer service, easy access to product information, and a smooth purchase process.</p> <p>The model's performance and predictions should be continuously monitored to assess their accuracy and impact.</p> <p>Monitor the effectiveness of targeted marketing campaigns and personalized offers. Evaluate how well the model's predictions align with actual purchase behavior and adjust strategies as needed.</p> <p>Continuously collect new data and feedback to refine the model. Periodically retrain the model with updated data to maintain accuracy and relevance.</p> <p>The model provides a data-driven approach to understanding customer behavior and market trends.</p>

	<p>Use the insights derived from the model for strategic planning and decision-making. This includes setting sales targets, budgeting for marketing efforts, and developing long-term strategies to attract and retain customers.</p> <p>The business can enhance its marketing strategies, optimize inventory management, and improve customer engagement. These actions will help drive growth, increase sales, and maximize ROI, ultimately leading to a more effective and data-driven approach to understanding and serving bike buyers.</p>

GitHub

Link GitHub - I am sharing my documents and file phyton.

	https://github.com/patriciarangelsantos/GDDA708_A2_WineQuality_BikeBuyersClean.git
--	---